

Any modern computer system will incorporate (at least) two levels of storage:

**primary storage:**

**random access memory (DRAM)**

typical capacity	4 GB to 32 GB
cost per GB	\$4.00
typical transfer rate	20,000 MB/sec

**secondary storage:**

**magnetic disk**

**SSD**

typical capacity	500 GB to 8 TB	128GB to 4TB
cost per GB	\$0.025	\$0.22
typical transfer rate	150 MB/sec	2000 MB/sec

**Note: all statistics here were obtained on Oct 22, 2020.**

## Spatial units:

byte (B)	8 bits
kibibyte (KiB)	1024 or $2^{10}$ bytes
mebibyte (MiB)	1024 kibibytes or $2^{20}$ bytes
gibibyte (GiB)	1024 mebibytes or $2^{30}$ bytes

### IEC standard

byte (B)	8 bits
kilobyte (KB)	1024 or $2^{10}$ bytes
megabyte (MB)	1024 kilobytes or $2^{20}$ bytes
gigabyte (GB)	1024 megabytes or $2^{30}$ bytes

### traditional

byte (B)	8 bits
kilobyte (KB)	1000 or $10^3$ bytes
megabyte (MB)	1000 kilobytes or $10^6$ bytes
gigabyte (GB)	1000 megabytes or $10^9$ bytes

### alt. industry

## Time units:

picosecond (ps)	one-trillionth ( $10^{-12}$ ) of a second
nanosecond (ns)	one-billionth ( $10^{-9}$ ) of a second
microsecond ( $\mu$ s)	one-millionth ( $10^{-6}$ ) of a second
millisecond (ms)	one-thousandth ( $10^{-3}$ ) of a second

While the particular values given earlier are volatile, the relative performances suggested are actually fairly stable over time:

Primary storage:

- costs 100-200 times as much per unit as secondary storage.
- has transfer rates that are perhaps 100-200 times faster

Why do WE care (in a data structures class)?

For many applications

- full data sets are too large to store in memory at once
- data must be first read from secondary storage into memory for processing
- and then results must be written back to secondary storage after processing

What can a programmer do to improve performance in disk-heavy applications?

- take an idea from the use of memory caches in hardware designs
- create an in-memory (DRAM) data structure to hold recently- and/or frequently-accessed records
- count on *locality of reference* in the application's record retrievals
- strive for the average record fetch to resemble a DRAM access rather than a secondary storage access

We call such a data structure a *buffer pool*.

In view of the previous discussion of secondary storage, it makes sense to design programs so that data is read from and written to disk in relatively large chunks... but there is more.

## *Temporal Locality of Reference*

In many cases, if a program accesses one part of a file, there is a high probability that the program will access the same part of the file again in the near future.

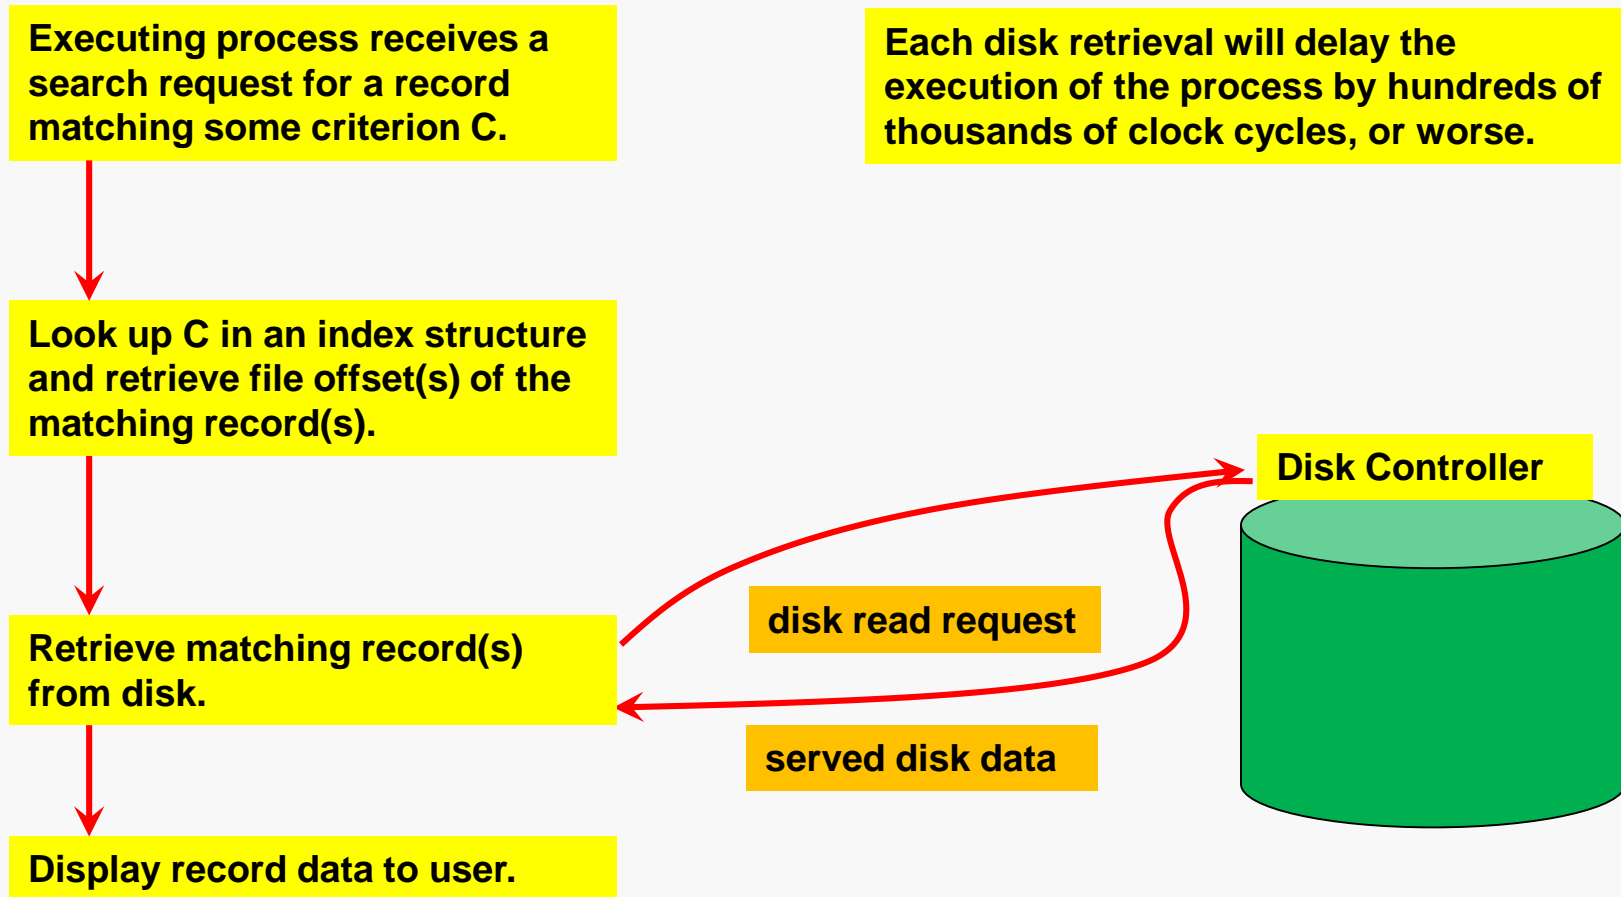
**Moral: once you've grabbed a chunk, keep it around.**

## *Spatial Locality of Reference*

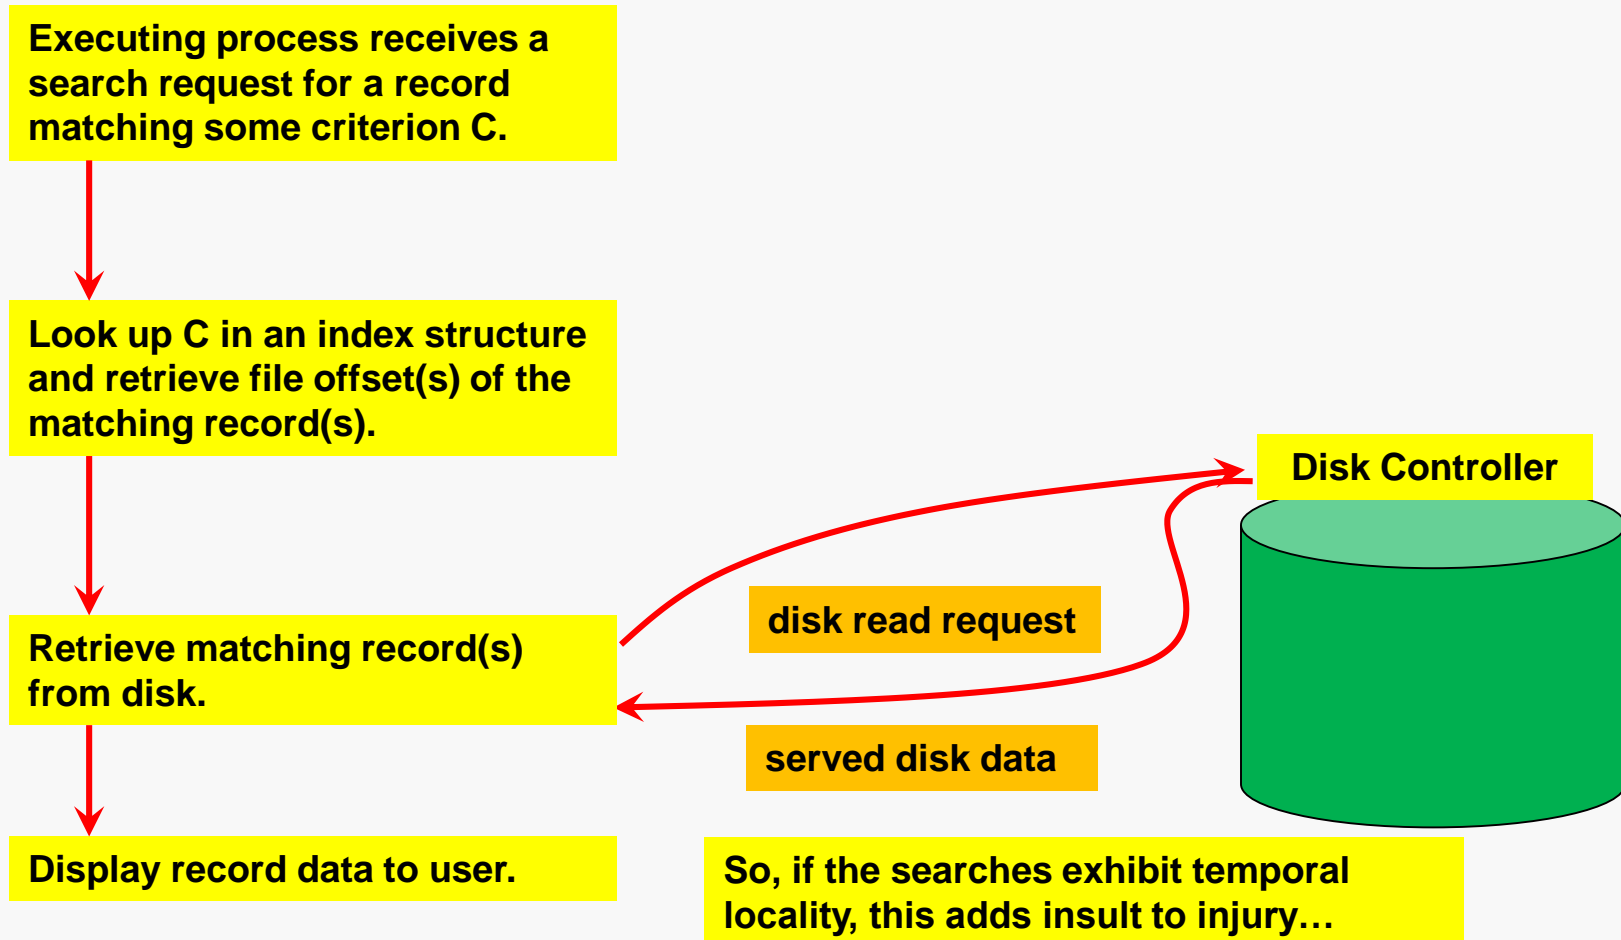
In many cases, if a program accesses one part of a file, there is a high probability that the program will access nearby parts of the file in the near future.

**Moral: grab a larger chunk than you immediately need.**

A program that retrieves records from disk in response to search requests would (naively) have interactions like this:



Not only does this hurt performance when a record is retrieved, we pay the same time cost if that same record is requested again...



*buffer pool* a series of buffers (memory locations) used by a program to cache disk data

Basically, the buffer pool is just a collection records, stored in RAM.

When a record is requested, the program first checks to see if the record is in the pool.

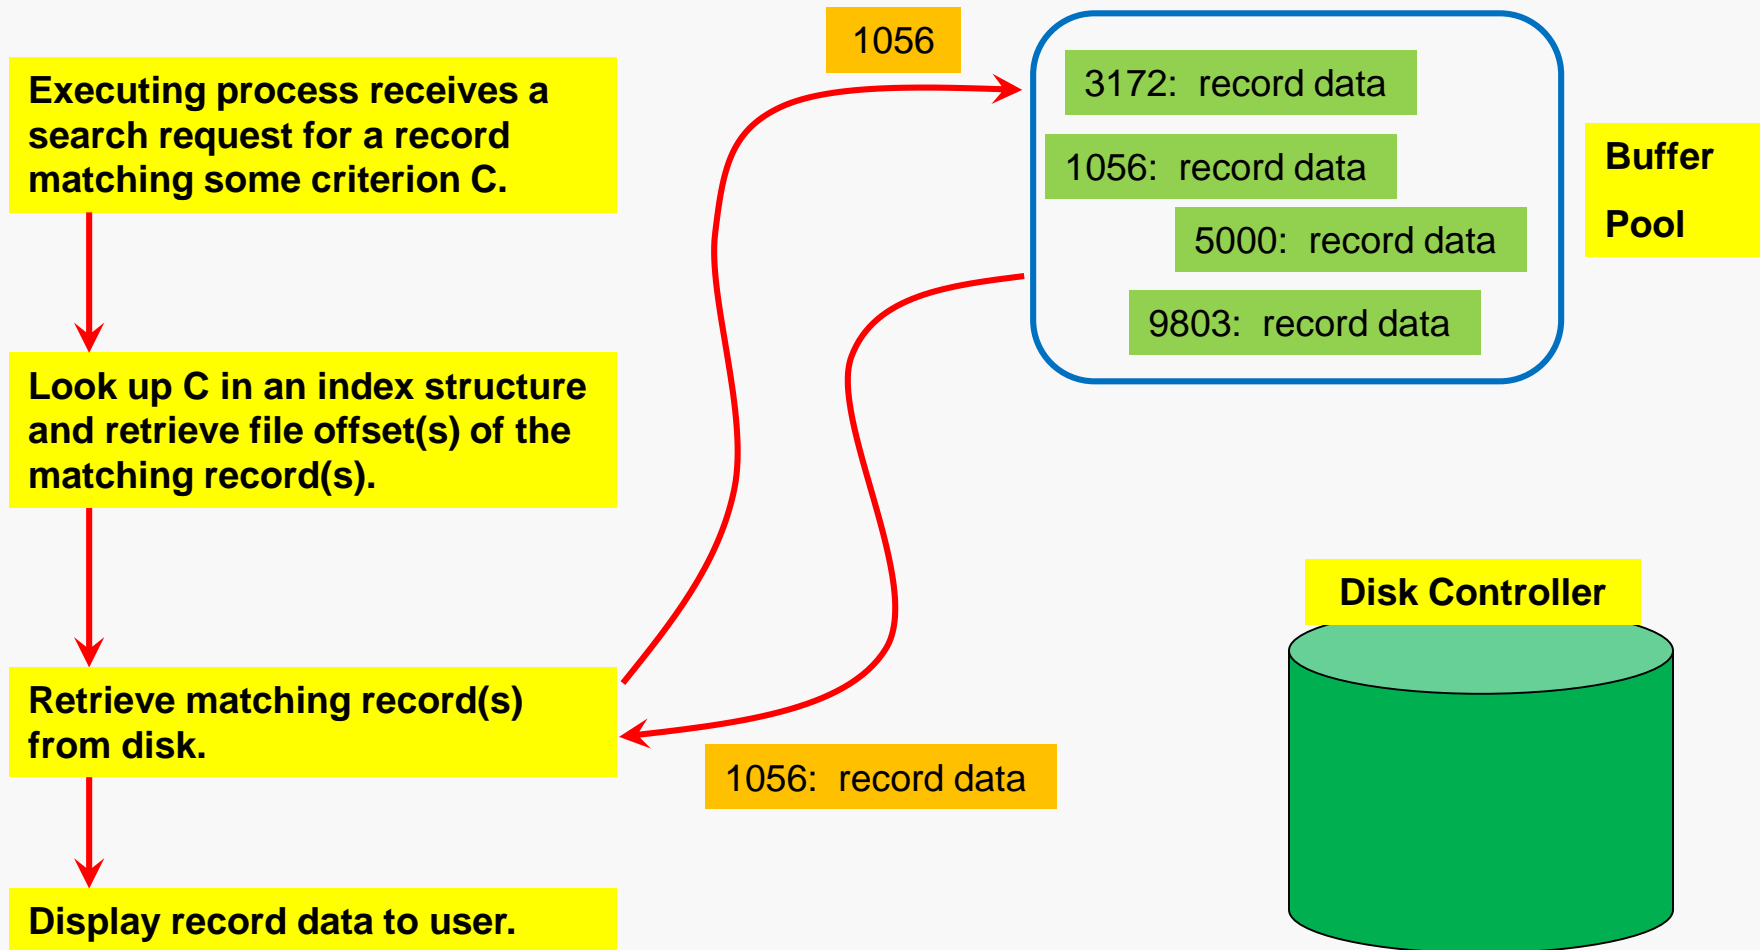
If so, there's no need to go to disk to get the record, and time is saved.

When the program does retrieve a record from disk, the newly-read record is copied into the pool, replacing a currently-stored record if necessary.



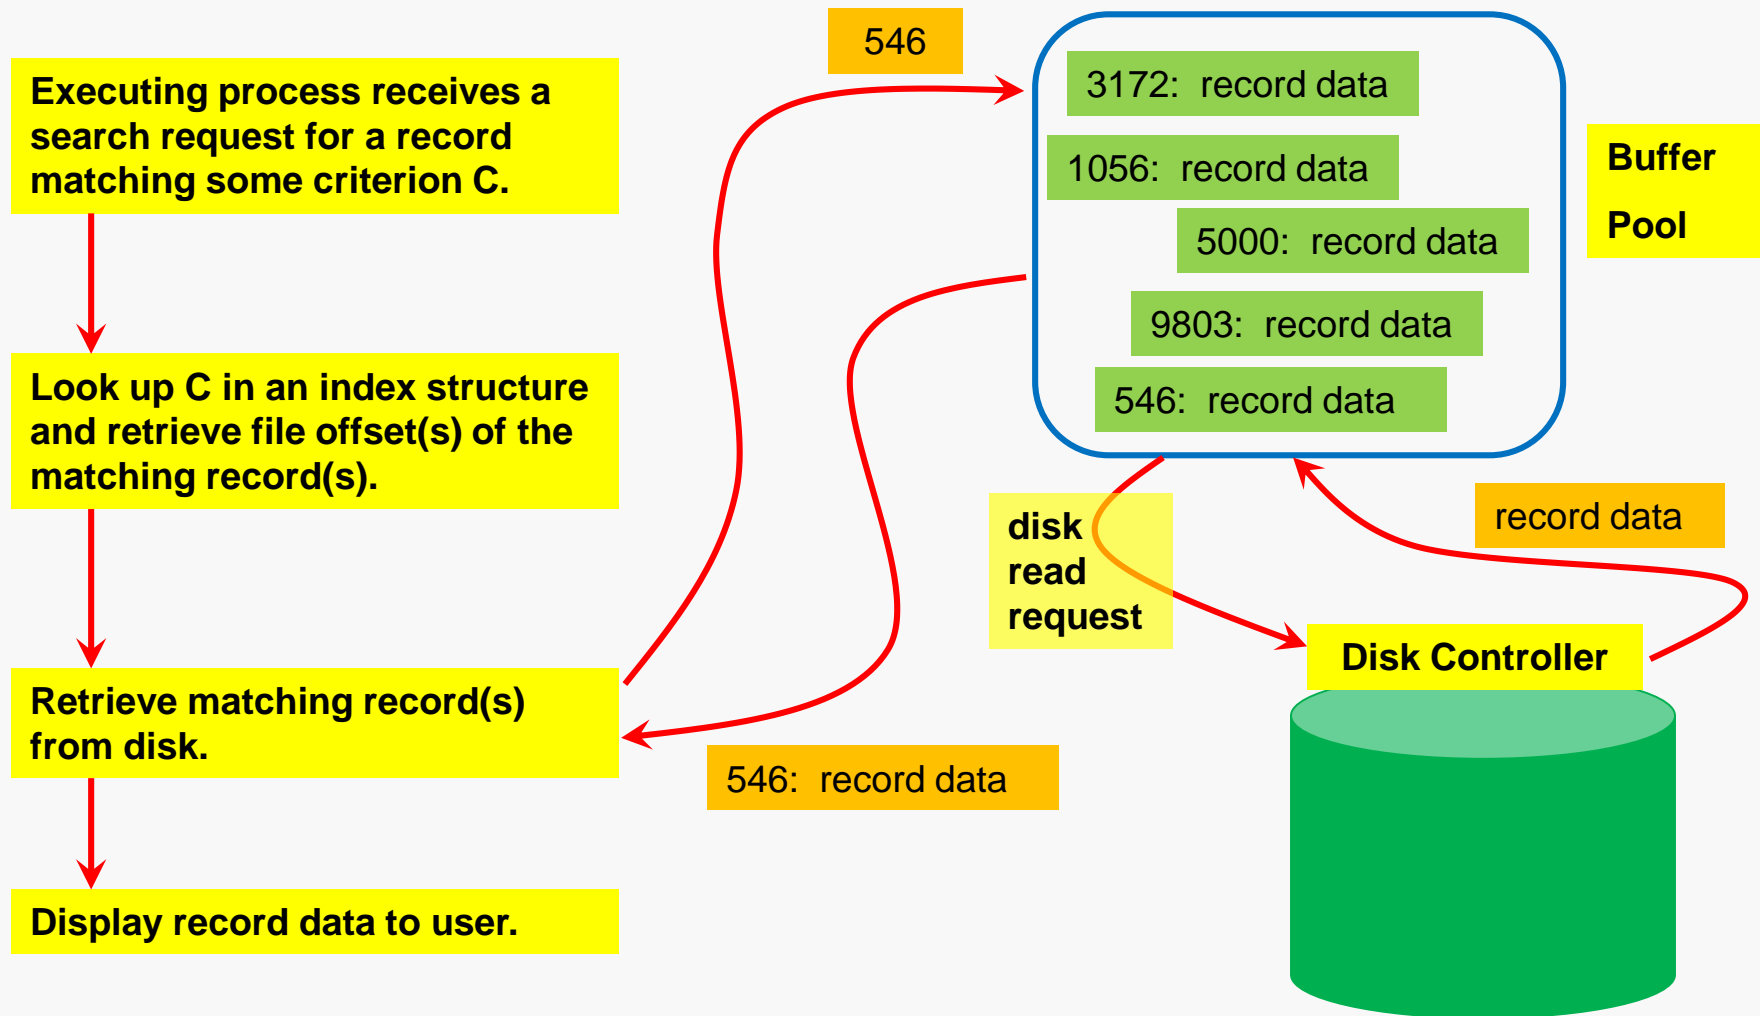
The interaction of the rest of the process with the disk is now mediated by the pool.

When we get a “hit”, we don’t go to disk:



The interaction of the rest of the process with the disk is now mediated by the pool.

When we get a “miss”, the pool goes to disk, updates itself, and serves up the record:



The buffer pool must be organized physically and logically.

The physical organization is generally an ordered list of some sort.

The logical organization depends upon how the buffer pool deals with the issue of replacement — if a new record must be added to the pool and all the buffers are currently full, one of the current records must be replaced.

If the replaced element has been modified, it (usually) must be written back to disk or the changes will be lost. Thus, some replacement strategies may include a consideration of which buffer elements have been modified in choosing one to replace.

Some common buffer replacement strategies:

FIFO (first-in is first-out) organize buffers as a queue

LFU (least frequently used) replace the least-accessed buffer

LRU (least recently used) replace the longest-idle buffer

Logically the buffer pool is treated as a queue:

```
655: 655 miss
289: 655 289 miss
586: 655 289 586 miss
289: 655 289 586 hit
694: 655 289 586 694 miss
586: 655 289 586 694 hit
655: 655 289 586 694 hit
138: 655 289 586 694 138 miss
289: 655 289 586 694 138 hit
694: 655 289 586 694 138 hit
289: 655 289 586 694 138 hit
694: 655 289 586 694 138 hit
851: 289 586 694 138 851 miss
586: 289 586 694 138 851 hit
330: 586 694 138 851 330 miss
289: 694 138 851 330 289 miss
694: 694 138 851 330 289 hit
331: 138 851 330 289 331 miss
289: 138 851 330 289 331 hit
694: 851 330 289 331 694 miss
Number of accesses: 20
Number of hits: 10
Number of misses: 10
Hit rate: 50.00
```

Takes no notice of the access pattern exhibited by the program. Consider what would happen with the sequence:

655

289

655

393

655

127

655

781

...

For LFU we must maintain an access count for each element of the buffer pool. It is also useful to keep the elements sorted by that count.

```

655:  (655, 1)    miss
289:  (655, 1)    (289, 1)    miss
586:  (655, 1)    (289, 1)    (586, 1)    miss
289:  (289, 2)    (655, 1)    (586, 1)    hit
694:  (289, 2)    (655, 1)    (586, 1)    (694, 1)    miss
586:  (289, 2)    (586, 2)    (655, 1)    (694, 1)    hit
655:  (289, 2)    (586, 2)    (655, 2)    (694, 1)    hit
138:  (289, 2)    (586, 2)    (655, 2)    (694, 1)    (138, 1)
289:  (289, 3)    (586, 2)    (655, 2)    (694, 1)    (138, 1)
694:  (289, 3)    (586, 2)    (655, 2)    (694, 2)    (138, 1)
289:  (289, 4)    (586, 2)    (655, 2)    (694, 2)    (138, 1)
694:  (289, 4)    (694, 3)    (586, 2)    (655, 2)    (138, 1)
851:  (289, 4)    (694, 3)    (586, 2)    (655, 2)    (851, 1)
586:  (289, 4)    (694, 3)    (586, 3)    (655, 2)    (851, 1)
330:  (289, 4)    (694, 3)    (586, 3)    (655, 2)    (330, 1)
289:  (289, 5)    (694, 3)    (586, 3)    (655, 2)    (330, 1)
694:  (289, 5)    (694, 4)    (586, 3)    (655, 2)    (330, 1)
331:  (289, 5)    (694, 4)    (586, 3)    (655, 2)    (331, 1)
289:  (289, 6)    (694, 4)    (586, 3)    (655, 2)    (331, 1)
694:  (289, 6)    (694, 5)    (586, 3)    (655, 2)    (331, 1)

Number of accesses: 20
Number of hits:     12
Number of misses:   8
Hit rate:           60.00
    
```

Aside from cost of storing and maintaining counter values, and searching for least value, consider the sequence:

655 (500 times)

289 (500 times)

100

101

102

103

...

With LRU, we may use a simple list structure. On an access, we move the targeted element to the front of the list. That puts the least recently used element at the tail of the list.

```

655: 655    miss
289: 289    655    miss
586: 586    289    655    miss
289: 289    586    655    hit
694: 694    289    586    655    miss
586: 586    694    289    655    hit
655: 655    586    694    289    hit
138: 138    655    586    694    289    miss
289: 289    138    655    586    694    hit
694: 694    289    138    655    586    hit
289: 289    694    138    655    586    hit
694: 694    289    138    655    586    hit
851: 851    694    289    138    655    miss
586: 586    851    694    289    138    miss
330: 330    586    851    694    289    miss
289: 289    330    586    851    694    hit
694: 694    289    330    586    851    hit
331: 331    694    289    330    586    miss
289: 289    331    694    330    586    hit
694: 694    289    331    330    586    hit
Number of accesses: 20
Number of hits:      11
Number of misses:    9
Hit rate:             55.00
    
```

Consider what would happen with the sequence:

655

289

655

301

302

303

304

289

...

You would (perhaps) expect that if you increased the number of slots in the pool, then for the same sequence of record references you'd get fewer misses (or at least not get more misses).

You may be disappointed, at least if you use FIFO replacement:

Record	Pool of size 3			
	X	X	X	
1	1	X	X	
2	1	2	X	
3	1	2	3	
4	2	3	4	
1	3	4	1	
2	4	1	2	
5	1	2	5	
1	1	2	5	hit!
2	1	2	5	hit!
3	2	5	3	
4	5	3	4	
5	5	3	4	hit!

Record	Pool of size 4				
	X	X	X	X	
1	1	X	X	X	
2	1	2	X	X	
3	1	2	3	X	
4	1	2	3	4	
1	1	2	3	4	hit!
2	1	2	3	4	hit!
5	2	3	4	5	
1	3	4	5	1	
2	4	5	1	2	
3	5	1	2	3	
4	1	2	3	4	
5	2	3	4	5	

L A Belady, R A Nelson, G S Shedler

*An anomaly in space-time characteristics of certain programs running in a paging machine*

CACM Volume 12, Issue 6, June 1969

The performance of a replacement strategy is commonly measured by its *fault rate*, i.e., the percentage of requests that require a new element to be loaded into the pool.

Some observations:

- misses will occur unless the pool contains the entire collection of data objects that are needed (the *working set*)
- which data objects are needed tends to change over time as the program runs, so the working set varies over time
- if the buffer pool is too small, it may be impossible to keep the current working set resident (in the buffer pool)
- if the buffer pool is too large, the program will waste memory



None of these replacement strategies, or any other feasible one, is best in all cases.

All are used with some frequency.

Intuitively, LRU and LFU make more sense than FIFO.

The performance you get is determined by the access pattern exhibited by the running program, and that is often impossible to predict.

Belady's optimal replacement strategy:

replace the element whose next access lies furthest in the future

Sometimes stated as “replace the element with the maximal forward distance”.

Requires knowing the future, and so is impossible to implement.

Does suggest considering predictive strategies.

Ideal replacement strategy:

Replace an element whose forward distance is maximal.

QTPs:

By what logic does FIFO estimate forward distance?

By what logic does LFU estimate forward distance?

By what logic does LRU estimate forward distance?

A buffer pool can service temporal locality:

Keep the fetched record around in RAM for awhile...

QTP: how could buffer pool service spatial locality?

There are some general properties a good buffer pool will have:

- the buffer size and number of buffers should be client-configurable
- the buffer pool may deal only in "raw bytes"; i.e., not know anything at all about the internals of the data record format used by the client code

OR

the buffer pool may deal in interpreted data records, parsed from the file and transformed into an object

- if records are fixed-length then each buffer should hold an integer number of records; for variable-length records, things are more complex and it is often necessary for buffers to allow some internal fragmentation
- empirically, a program using a buffer pool is considered to be achieving good performance if less than 10% of the record references require loading a new record into the buffer pool