

*hash function* a function that can take a key value and compute an integer value (or an index in a table) from it

For example, student records for a class could be stored in an array  $C$  of dimension 10000 by truncating the student's ID number to its last four digits:

$$H(\text{IDNum}) = \text{IDNum} \% 10000$$

Given an ID number  $X$ , the corresponding record would be inserted at  $C[H(X)]$ .

This would be easy to implement, and cheap to execute. Whether it's actually a very good hash function is another matter...

Suppose we have  $N$  records, and a table of  $M$  slots, where  $N \leq M$ .

- there are  $M^N$  different ways to map the records into the table, if we don't worry about mapping two records to the same slot
- the number of different *perfect* mappings of the records into different slots in the table would be


$$P(M, N) = \frac{M!}{(M - N)!}$$

- for instance, if  $N = 50$  and  $M = 100$ , there are  $100^{50} = 10^{100}$  different possible hash mappings, “only” about  $10^{93}$  of which are perfect (1 in 10,000,000)
- so, there is no shortage of potential perfect hash functions (in theory)
- however, we need one that is effectively computable, that is, it must be possible to compute it (so we need a formula for it) and it must be efficiently computable
- there are a number of common approaches, but the design of good, practical hash functions must still be considered a topic of research and experiment

The set of logically possible key values may be very large.

- set of possible Java identifiers of length 10 or less (xxx)


The set of key values we actually encounter when compiling a program will be much smaller, but we don't know which values we'll actually see until we see them...



The ideal is a one-to-one hash function... good luck with that:

- take a reasonable table size for hashing the identifiers in a Java program
- consider the number of possible Java identifiers
- both sets are finite and the second is much, much larger

**F() may be uniform on the whole theoretical domain...**



So, the next best thing would be a hash function that is "uniform".

That is, we'd like to map about the same number of domain values to each slot in the table... good luck with that too...

**...but not at all uniform on the small subset of it that we actually get.**

It is usually desirable to have the entire key value affect the hash result (so simply chopping off the last k digits of an integer key is NOT a good idea in most cases).

Consider the following function to hash a string value into an integer range:

```
public static int sumOfChars(String toHash) {  
  
    int hashValue = 0;  
    for (int Pos = 0; Pos < toHash.length(); Pos++) {  
        hashValue = hashValue + toHash.charAt(Pos);  
    }  
    return hashValue;  
}
```

**Hashing: hash**  
h: 104  
a: 97  
s: 115  
h: 104  
Sum: 420  
  
Mod by table  
size to get the  
index

This takes every element of the string into account... a string hash function that truncated to the last three characters would compute the same integer for "hash", "stash", "mash", "trash".

## Division

- the first order of business for a hash function is to compute an integer value
- if we expect the hash function to produce a valid index for our chosen table size, that integer will probably be out of range
- that is easily remedied by modding the integer by the table size
- perhaps it is better if the table size is a prime, or at least has no small prime factors

## Folding

- portions of the key are often recombined, or *folded* together
- *shift folding*:  $123-45-6789 \rightarrow 123 + 456 + 789$
- *boundary folding*:  $123-45-6789 \rightarrow 123 + 654 + 789$
- can be efficiently performed using bitwise operations
- the characters of a string can be xor'd together, but small numbers result
- “chunks” of characters can be xor'd instead, say in integer-sized chunks

## Mid-square function

- square the key, then use the middle part as the result
- e.g.,  $3121 \rightarrow 9740641 \rightarrow 406$  (with a table size of 1000)
- a string would first be transformed into a number, say by folding
- idea is to let all of the key influence the result
- if table size is a power of 2, this can be done efficiently at the bit level:  
 $3121 \rightarrow 100101001010000101100001 \rightarrow 0101000010$  (with a table size of 1024)

## Extraction

- use only part of the key to compute the result
- motivation may be related to the distribution of the actual key values, e.g., VT student IDs almost all begin with 904, so it would contribute no useful separation

A good hash function should:

- be easy and quick to compute
- achieve an even distribution of the key values that actually occur across the index range supported by the table
- ideally be mathematically one-to-one on the set of relevant key values

Note: hash functions are NOT random in any sense.



A simple hash function is likely to map two or more key values to the same integer value, in at least some cases.

A little bit of design forethought can often reduce this:

```
public static int sumOfShiftedChars(String toHash) {  
  
    int hashValue = 0;  
    for (int Pos = 0; Pos < toHash.length(); Pos++) {  
        hashValue = (hashValue << 4) + toHash.charAt(Pos);  
    }  
    return hashValue;  
}
```

**Hashing: hash**

h: 104

a: 97

s: 115

h: 104

**Sum: 452760**

The original version would have hashed both of these strings to the same table index.

Flaw: it didn't take element position into account.

**Hashing: shah**

s: 115

h: 104

a: 97

h: 104

**Sum: 499320**

Consider the following function to hash a string value into an integer:

```
public static long elfHash(String toHash) {  
  
    long hashValue = 0;  
    for (int Pos = 0; Pos < toHash.length(); Pos++) {           // use all elements  
  
        hashValue = (hashValue << 4) + toHash.charAt(Pos);      // shift/mix  
  
        long hiBits = hashValue & 0xF000000000000000L;          // get high nybble  
  
        if (hiBits != 0)  
            hashValue ^= hiBits >> 56;                          // xor high nybble with second nybble  
  
        hashValue &= ~hiBits;                                    // clear high nybble  
    }  
  
    return hashValue;  
}
```

This was developed originally during the design of the UNIX operating system, for use in building system-level hash tables.

Here's a trace (using 32-bit integers):

Character	hashValue
-----	
d: 64	00000064
i: 69	000006a9
s: 73	00006b03
t: 74	0006b0a4
r: 72	006b0ab2
i: 69	06b0ab89
b: 62	0b0ab892
u: 75	00ab8925
t: 74	0ab892c4
i: 69	0b892c09
o: 6f	0892c04f
n: 6e	092c05de
distribution: 15388030	

```

hashValue      : 06b0ab89
hashValue << 4 : 6b0ab890
add 62         : 6b0ab8f2

hiBits         : 60000000
hiBits >> 24   : 00000060

hashValue ^    6b0ab8f2
  hiBits      00000060
              : 6b0ab892

hashValue &
  ~hiBits     : 0b0ab892
    
```

```

f: 1111
6: 0110
^: 1001
    
```